

Dr. Sarah Lin

AI / Machine Learning Engineer | sarah.lin@example.com | Seattle, WA

PROFESSIONAL SUMMARY

Machine Learning and AI Systems Engineer with 4+ years of experience designing deep learning models, RAG agentic pipelines, and MLOps deployment architectures. Specialized in PyTorch, TensorFlow, LLMs, LangChain, Vector Databases, and production model serving.

TECHNICAL SKILLS & COMPETENCIES

ML & Deep Learning: Python, PyTorch, TensorFlow, Scikit-Learn, Keras, HuggingFace, Transformers

Generative AI & LLMs: LLMs, LangChain, LlamaIndex, RAG Pipelines, Vector Databases (Pinecone, ChromaDB), Prompt Engineering

MLOps & Cloud: Docker, Kubernetes, MLflow, Triton Inference Server, AWS, GCP, Git, CI/CD

Data & Architecture: SQL, Pandas, NumPy, FastAPI, Distributed Training, Model Quantization

WORK EXPERIENCE

Senior Machine Learning Engineer — *Synthetix AI Lab* (2022 - Present)

- Fine-tuned open-source LLMs using PyTorch, HuggingFace, and LoRA/QLoRA for enterprise question-answering systems.
- Architected high-throughput RAG pipelines using LangChain, Vector Databases, and FastAPI, serving 10,000+ daily user requests.
- Implemented automated model evaluation, tracking, and deployment using MLflow and Docker containers on Kubernetes.

FEATURED PROJECTS

Multi-Agent RAG Knowledge Copilot | *Tech: Python, PyTorch, LangChain, Vector Databases, FastAPI, Docker*

- Engineered an autonomous multi-agent document analysis system utilizing semantic embeddings and hybrid search retrieval.

Distributed Computer Vision & Inference Pipeline | *Tech: Python, TensorFlow, PyTorch, Triton, Docker, Git*

- Deployed optimized deep learning vision models with ONNX Runtime, achieving 4x throughput improvement on GPU nodes.

EDUCATION & CERTIFICATIONS

- Ph.D. in Artificial Intelligence & Computer Science — University of Washington
- AWS Certified Machine Learning - Specialty